

```
C:\HADOOPOUTPUT\spark>spark-submit --verbose wordcountSpark.jar -class JavaWordCount yarn-client
```

The master URL passed to Spark can be in one of the following formats:

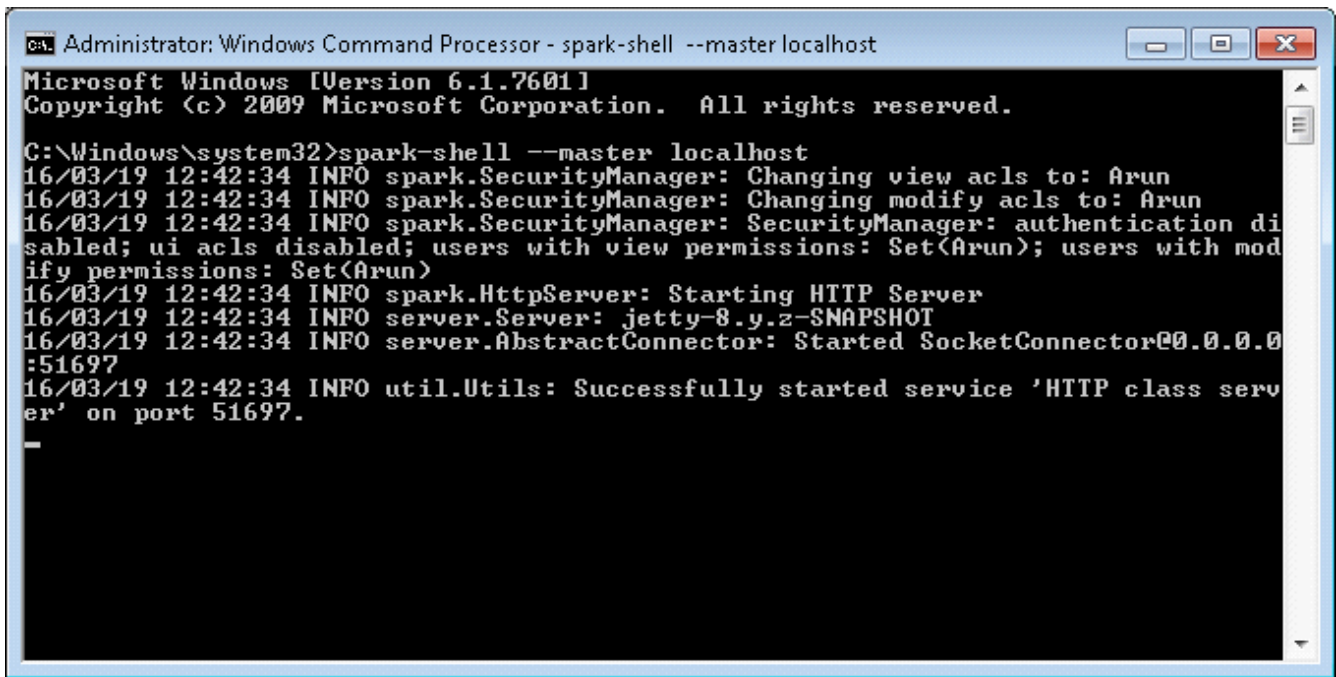
Master URL	Meaning
local	Run Spark locally with one worker thread (i.e. no parallelism at all).
local[K]	Run Spark locally with K worker threads (ideally, set this to the number of cores on your machine).
local[*]	Run Spark locally with as many worker threads as logical cores on your machine.
spark://HOST:PORT	Connect to the given Spark standalone cluster master. The port must be whichever one your master is configured to use, which is 7077 by default.
mesos://HOST:PORT	Connect to the given Mesos cluster. The port must be whichever one your is configured to use, which is 5050 by default. Or, for a Mesos cluster using ZooKeeper, use <code>mesos://zk://...</code> . To submit with <code>--deploy-mode cluster</code> , the HOST:PORT should be configured to connect to the MesosClusterDispatcher .
yarn	Connect to a YARN cluster in <code>client</code> or <code>cluster</code> mode depending on the value of <code>--deploy-mode</code> . The cluster location will be found based on the <code>HADOOP_CONF_DIR</code> or <code>YARN_CONF_DIR</code> variable.
yarn-client	Equivalent to <code>yarn</code> with <code>--deploy-mode client</code> , which is preferred to <code>'yarn-client'</code>
yarn-cluster	Equivalent to <code>yarn</code> with <code>--deploy-mode cluster</code> , which is preferred to <code>'yarn-cluster'</code>

- **file:** - Absolute paths and `file:/` URIs are served by the driver's HTTP file server, and every executor pulls the file from the driver HTTP server.
- **hdfs:, http:, https:, ftp:** - these pull down files and JARs from the URI as expected
- **local:** - a URI starting with `local:/` is expected to exist as a local file on each worker node. This means that no network IO will be incurred, and works well for large files/JARs that are pushed to each worker, or shared via NFS, GlusterFS, etc.

NOTE:

To avoid block in use or hdfs running in safe mode block every time starting hadoop, set **YARN_CONFIG_DIR**

```
YARN_CONFIG_DIR-C:\hadoop-2.2.0\etc\hadoop
```



```
Administrator: Windows Command Processor - spark-shell --master localhost
Microsoft Windows [Version 6.1.7601]
Copyright (c) 2009 Microsoft Corporation. All rights reserved.

C:\Windows\system32>spark-shell --master localhost
16/03/19 12:42:34 INFO spark.SecurityManager: Changing view acls to: Arun
16/03/19 12:42:34 INFO spark.SecurityManager: Changing modify acls to: Arun
16/03/19 12:42:34 INFO spark.SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view permissions: Set(Arun); users with modify permissions: Set(Arun)
16/03/19 12:42:34 INFO spark.HttpServer: Starting HTTP Server
16/03/19 12:42:34 INFO server.Server: jetty-8.y.z-SNAPSHOT
16/03/19 12:42:34 INFO server.AbstractConnector: Started SocketConnector@0.0.0.0:51697
16/03/19 12:42:34 INFO util.Utils: Successfully started service 'HTTP class server' on port 51697.
```

above fails finally in localhost
start with spark-shell alone.

>spark-shell

Spark copies the Spark assembly JAR file to HDFS each time you run spark-submit. You can avoid doing this copy each time by manually uploading the Spark assembly JAR file to your HDFS. Then, set the SPARK_JAR environment variable to this HDFS path

```
C:\spark-1.6.0-bin-hadoop2.3\lib>hdfs dfs -put ./spark-assembly-1.6.0-hadoop2.3.0.jar /spark/lib/spark-assembly.jar
```

Reading data from hdfs

Files in hdfs are usually stored in the following formats:

- plain txt/csv/json files
- [sequence files](#). You can think of them as serialized java objects. In recent years became less popular. Also they are not portable (need custom readers), so I do not find them interesting for this post.
- [avro](#)(row-based)
- [parquet](#)(column-based)
- [orc](#)(column-based)

Good news is that Spark (and SparkR!) can read `json`, `parquet`, `orc` with built-in `read.df` function and `csv`, `avro` with `read.df` and `spark-avro`, `spark-csv` spark packages.

```
Administrator: Windows Command Processor
C:\spark-1.6.0-bin-hadoop2.3\lib>hdfs dfs -ls /spark/*
C:\spark-1.6.0-bin-hadoop2.3\lib>hdfs dfs -ls /spark/
Found 1 items
drwxr-xr-x - Arun supergroup 0 2016-03-24 03:47 /spark/lib
C:\spark-1.6.0-bin-hadoop2.3\lib>hdfs dfs -put C:\spark-1.6.0-bin-hadoop2.3\lib\
spark-assembly-1.6.0-hadoop2.3.0.jar \spark\lib\spark-assembly-1.6.0-hadoop2.3
.0.jar
put: unexpected URISyntaxException
put: unexpected URISyntaxException
put: 'sparklibspark-assembly-1.6.0-hadoop2.3.0.jar': No such file or directory
C:\spark-1.6.0-bin-hadoop2.3\lib>hdfs dfs -put spark-assembly-1.6.0-hadoop2.3.0.
jar \spark\lib\spark-assembly.jar
put: 'sparklibspark-assembly.jar': No such file or directory
C:\spark-1.6.0-bin-hadoop2.3\lib>hdfs dfs -put ./spark-assembly-1.6.0-hadoop2.3.
0.jar \spark\lib\spark-assembly.jar
put: 'sparklibspark-assembly.jar': No such file or directory
C:\spark-1.6.0-bin-hadoop2.3\lib>hdfs dfs -put ./spark-assembly-1.6.0-hadoop2.3.
0.jar /spark/lib/spark-assembly.jar
C:\spark-1.6.0-bin-hadoop2.3\lib>
```

NOTE:

for running java use cmd run-example

For Scala and Java, use run-example:

```
./bin/run-example SparkPi
```

For Python examples, use spark-submit directly:

```
./bin/spark-submit examples/src/main/python/pi.py
```

For R examples, use spark-submit directly:

```
./bin/spark-submit examples/src/main/r/dataframe.R
```

RUNNING JAVA IN SPARK

```
public class SparkExample {
    public static final int NUM_SAMPLES=10;
    SparkExample(){
        SparkConf conf = new SparkConf().setAppName("Spark Pi");
        conf.setMaster("local[2]");
        SparkContext sc = new SparkContext(conf);

        RDD<String> textFile =sc.textFile("hdfs://input/wordcount.txt",1);
        System.out.println("SparkExample.SparkExample()"+textFile);
    }
}
```

```

    }

    public static void main(String args[]){
        new SparkExample();
    }
}

```

OUTPUT:

```

16/03/24 05:39:55 INFO BlockManagerMasterEndpoint: Registering block manager
localhost:56793 with 795.0 MB RAM, BlockManagerId(driver, localhost, 56793)
16/03/24 05:39:55 INFO BlockManagerMaster: Registered BlockManager
16/03/24 05:39:56 INFO MemoryStore: Block broadcast_0 stored as values in memory
(estimated size 104.0 KB, free 104.0 KB)
16/03/24 05:39:57 INFO MemoryStore: Block broadcast_0_piece0 stored as bytes in memory
(estimated size 9.8 KB, free 113.8 KB)
16/03/24 05:39:57 INFO BlockManagerInfo: Added broadcast_0_piece0 in memory on
localhost:56793 (size: 9.8 KB, free: 795.0 MB)
16/03/24 05:39:57 INFO SparkContext: Created broadcast 0 from textFile at
SparkExample.java:15
SparkExample.SparkExample()hdfs://input/wordcount.txt MapPartitionsRDD[1] at textFile at
SparkExample.java:15
16/03/24 05:39:57 INFO SparkContext: Invoking stop() from shutdown hook
16/03/24 05:39:57 INFO SparkUI: Stopped Spark web UI at http://localhost:4041
16/03/24 05:39:57 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint st

```

MAP REDUCE IN SPARK WITHOUT USING HADOOP

```

package spark;

import java.util.Arrays;
import java.util.List;
import java.util.regex.Pattern;

import org.apache.spark.SparkConf;
import org.apache.spark.api.java.JavaPairRDD;
import org.apache.spark.api.java.JavaRDD;
import org.apache.spark.api.java.JavaSparkContext;
import org.apache.spark.api.java.function.FlatMapFunction;
import org.apache.spark.api.java.function.Function2;
import org.apache.spark.api.java.function.PairFunction;

import scala.Tuple2;

public final class JavaWordCount {
    private static final Pattern SPACE = Pattern.compile(" ");

    public static void main(String[] args) throws Exception {

        SparkConf sparkConf = new SparkConf().setAppName("JavaWordCount");
        sparkConf.setMaster("local[1]");
    }
}

```

```

    JavaSparkContext ctx = new JavaSparkContext(sparkConf);
    JavaRDD<String> lines = ctx.textFile("c:/HADOOPOUTPUT/wordcount.txt", 1);
    System.out.println("JavaWordCount.main()"+lines);
    JavaRDD<String> words = lines.flatMap(new FlatMapFunction<String, String>() {
        @Override
        public Iterable<String> call(String s) {
            return Arrays.asList(SPACE.split(s));
        }
    });

    JavaPairRDD<String, Integer> ones = words.mapToPair(new PairFunction<String, String,
Integer>() {
        @Override
        public Tuple2<String, Integer> call(String s) {
            return new Tuple2<String, Integer>(s, 1);
        }
    });

    JavaPairRDD<String, Integer> counts = ones.reduceByKey(new Function2<Integer, Integer,
Integer>() {
        @Override
        public Integer call(Integer i1, Integer i2) {
            return i1 + i2;
        }
    });

    List<Tuple2<String, Integer>> output = counts.collect();
    for (Tuple2<?,?> tuple : output) {
        System.out.println(tuple._1() + ": " + tuple._2());
    }
    ctx.stop();
}
}

```

```

16/03/24 06:04:11 INFO TaskSchedulerImpl: Removed TaskSet 1.0, whose tasks have all
completed, from pool
16/03/24 06:04:11 INFO DAGScheduler: Job 0 finished: collect at JavaWordCount.java:67,
took 0.650911 s
example: 1
are: 1
is: 1
you: 1
wordcount: 1
hadoop: 3
hi: 3
how: 2
16/03/24 06:04:11 INFO SparkUI: Stopped Spark web UI at http://localhost:4041
16/03/24 06:04:11 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint
stopped!

```

See below same program run in hadoop takes 23 sec

```

        if(args!=null && args.length<1){
            args= new String[2];
            args[0]="c:/HADOOPOUTPUT/wordcount.txt";
            args[1]="c:/HADOOPOUTPUT/output";
        }
        Path inputPath = new Path(args[0]);
        Path outputPath = new Path(args[1]);
        Configuration conf = getConf();
        Job job = new Job(conf, this.getClass().toString());
        FileInputFormat.setInputPaths(job, inputPath);
        FileOutputFormat.setOutputPath(job, outputPath);
        job.setJobName("WordCount");
        job.setJarByClass(WordCount.class);
        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);
        job.setMapOutputKeyClass(Text.class);
        job.setMapOutputValueClass(IntWritable.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        job.setMapperClass(Map.class);
        job.setCombinerClass(Reduce.class);
        job.setReducerClass(Reduce.class);
        return job.waitForCompletion(true) ? 0 : 1;
    }

```

For Record size 1000:

```

16/03/24 06:27:40 INFO Executor: Running task 0.0 in stage 1.0 (TID 1)
16/03/24 06:27:40 INFO ShuffleBlockFetcherIterator: Getting 1 non-empty blocks out of 1
blocks
16/03/24 06:27:40 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 7 ms
16/03/24 06:27:41 INFO Executor: Finished task 0.0 in stage 1.0 (TID 1). 243558 bytes
result sent to driver
16/03/24 06:27:41 INFO TaskSetManager: Finished task 0.0 in stage 1.0 (TID 1) in 776 ms on
localhost (1/1)
16/03/24 06:27:41 INFO DAGScheduler: ResultStage 1 (collect at JavaWordCount.java:67)
finished in 0.777 s
16/03/24 06:27:41 INFO TaskSchedulerImpl: Removed TaskSet 1.0, whose tasks have all
completed, from pool
16/03/24 06:27:41 INFO DAGScheduler: Job 0 finished: collect at JavaWordCount.java:67,
took 2.073223 s
10:47,Product1,1200,Mastercard,David,Stavenger,Rogaland,country898,1/13/2009: 1
(Bruxelles),country574,1/31/2008: 1
,,,,,,country1956: 1
,,,,,,country2735: 1
8:23,Product1,1200,Mastercard,Sarah,Floyds: 1
House: 1
,CA,country82,1/3/2009: 1
22:00,Product2,3600,Amex,Lucien,Wiesbaden,Hessen,country233,1/4/2009: 1
,,,,,,country2781: 1

```

,, , , , , , C

PERFORMANCE COMPARISION

Comparison of time taken in diff environment.This proves hadoop is good for running small data also

Data limit	Standaone Java	Hadoop	Spark
5 lines of word count	2sec	23sec	0.6 sec(0.79 with hdfs)
1000 lines of csv	2 sec	22 sec	1.52 sec
3500 lines of csv	3sec	28 sec	2.313126 sec
10,000 lines of csv	4sec	26 sec	2.055446sec
65,536 lines of csv(max csv limit)	4sec	26-27	2.998 sec

ACCESS HDFS USING SPARK -JAVA

Spark can accesss hdfs file using hadoop-common-2.2.0.jar(version of inslalled hadoop) in classpath.you need to access using hdfs://localhost:9000/input/wordcount.txt.
/input/wordcount.txt is the hdfs file path.

```
package spark;

import java.util.Arrays;
import java.util.List;
import java.util.regex.Pattern;

import org.apache.hadoop.fs.Hdfs;
import org.apache.hadoop.io.BytesWritable;
import org.apache.hadoop.io.IntWritable;
import org.apache.spark.SparkConf;
import org.apache.spark.api.java.JavaPairRDD;
import org.apache.spark.api.java.JavaRDD;
import org.apache.spark.api.java.JavaSparkContext;
import org.apache.spark.api.java.function.FlatMapFunction;
import org.apache.spark.api.java.function.Function2;
import org.apache.spark.api.java.function.PairFunction;

import scala.Tuple2;

public final class JavaWordCount {
    private static final Pattern SPACE = Pattern.compile(" ");

    public static void main(String[] args) throws Exception {
```

```

SparkConf sparkConf = new SparkConf().setAppName("JavaWordCount");
sparkConf.setMaster("local[*]");
JavaSparkContext ctx = new JavaSparkContext(sparkConf);
// JavaRDD<String> lines = ctx.textFile("C:/HADOOPOUTPUT/SalesJan10000.csv", 1);

JavaRDD<String> lines = ctx.textFile("hdfs://localhost:9000/input/wordcount.txt", 1);
System.out.println("JavaWordCount.main()" + lines);
JavaRDD<String> words = lines.flatMap(new FlatMapFunction<String, String>() {
    @Override
    public Iterable<String> call(String s) {
        return Arrays.asList(SPACE.split(s));
    }
});

JavaPairRDD<String, Integer> ones = words.mapToPair(new PairFunction<String, String,
Integer>() {
    @Override
    public Tuple2<String, Integer> call(String s) {
        return new Tuple2<String, Integer>(s, 1);
    }
});

JavaPairRDD<String, Integer> counts = ones.reduceByKey(new Function2<Integer, Integer,
Integer>() {
    @Override
    public Integer call(Integer i1, Integer i2) {
        return i1 + i2;
    }
});

List<Tuple2<String, Integer>> output = counts.collect();
for (Tuple2<?, ?> tuple : output) {
    System.out.println(tuple._1() + ": " + tuple._2());
}
ctx.stop();
}
}

```

O/P;

```

example: 1
are: 1
is: 1
you: 1
wordcount: 1
hadoop: 3
hi: 3
how: 2

```

16/03/30 01:24:43 INFO DAGScheduler: Job 0 finished: collect at JavaWordCount.java:76, took 0.790907 s

NOTE:

1) Access hdfs using `hdfs://localhost:<port>/<hdfs file input>`

2) for spark with 5 lines of word count program without hdfs it uses **0.6sec** when access from local file system but when using hdfs it takes **0.79**.

As you can see it takes little longer when using hdfs for small size files.

3) hdfs files shown below

```
C:\Windows\system32>hdfs dfs -ls /input/*
```

Found 1 items

```
-rw-r--r-- 1 Arun supergroup 123637 2016-02-24 02:11 /input/sales.csv
```

Found 1 items

```
-rw-r--r-- 1 Arun supergroup 1398907 2016-02-25 00:09 /input/sales10000.csv
```

Found 1 items

```
-rw-r--r-- 1 Arun supergroup 466379 2016-02-24 22:53 /input/sales3500.csv
```

Found 1 items

```
-rw-r--r-- 1 Arun supergroup 8594762 2016-02-25 00:22 /input/sales65536.csv
```

Found 1 items

```
-rw-r--r-- 1 Arun supergroup 129745 2016-03-03 01:29 /input/salesunique.csv
```

v

Found 1 items

```
-rw-r--r-- 1 Arun supergroup 179820 2016-03-03 01:57 /input/salesunique3500.csv
```

Found 1 items

```
-rw-r--r-- 1 Arun supergroup 1476056 2016-03-03 01:47 /input/salesunique65536.csv
```

Found 1 items

```
-rw-r--r-- 1 Arun supergroup 70 2016-02-24 02:11 /input/wordcount.txt
```

RUN SPARK FROM CMD PROMPT -JAVA

```
C:\HADOOP\OUTPUT\spark>spark-submit --verbose wordcountSpark1.jar -class spark.avaWordCount
```

```
Administrator: Windows Command Processor
<TID 1> in 131 ms on localhost <1/1>
16/03/30 22:30:48 INFO scheduler.TaskSchedulerImpl: Removed TaskSet 1.0, whose t
asks have all completed, from pool
16/03/30 22:30:48 INFO scheduler.DAGScheduler: ResultStage 1 (collect at JavaWo
rdCount.java:76) finished in 0.132 s
16/03/30 22:30:48 INFO scheduler.DAGScheduler: Job 0 finished: collect at JavaWo
rdCount.java:76, took 2.372083 s
example: 1
are: 1
is: 1
you: 1
wordcount: 1
hadoop: 3
hi: 3
how: 2
16/03/30 22:30:48 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHan
dler{/metrics/json,null}
16/03/30 22:30:48 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHan
dler{/stages/stage/kill,null}
16/03/30 22:30:48 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHan
dler{/api,null}
16/03/30 22:30:48 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHan
dler{/,null}
16/03/30 22:30:48 INFO handler.ContextHandler: stopped o.s.j.s.ServletContextHan
dler{/static,null}
```

ACCESS HDFS USING SPARK -R Programming

Reading data from hdfs

Files in hdfs are usually stored in the following formats:

- plain txt/csv/json files
- [sequence files](#). You can think of them as serialized java objects. In recent years became less popular. Also they are not portable (need custom readers), so I do not find them interesting for this post.
- [avro](#)(row-based)
- [parquet](#)(column-based)
- [orc](#)(column-based)

Good news is that Spark (and SparkR!) can read `json`, `parquet`, `orc` with built-in `read.df` function and `csv`, `avro` with `read.df` and `spark-avro`, `spark-csv` spark packages.

STEP 1:

install devtools

STEP 2:

install sparkR

STEP 3:

Two ways to install Spark R:

Procedure 1:

when command below is given it automatically download Rbuild Tools

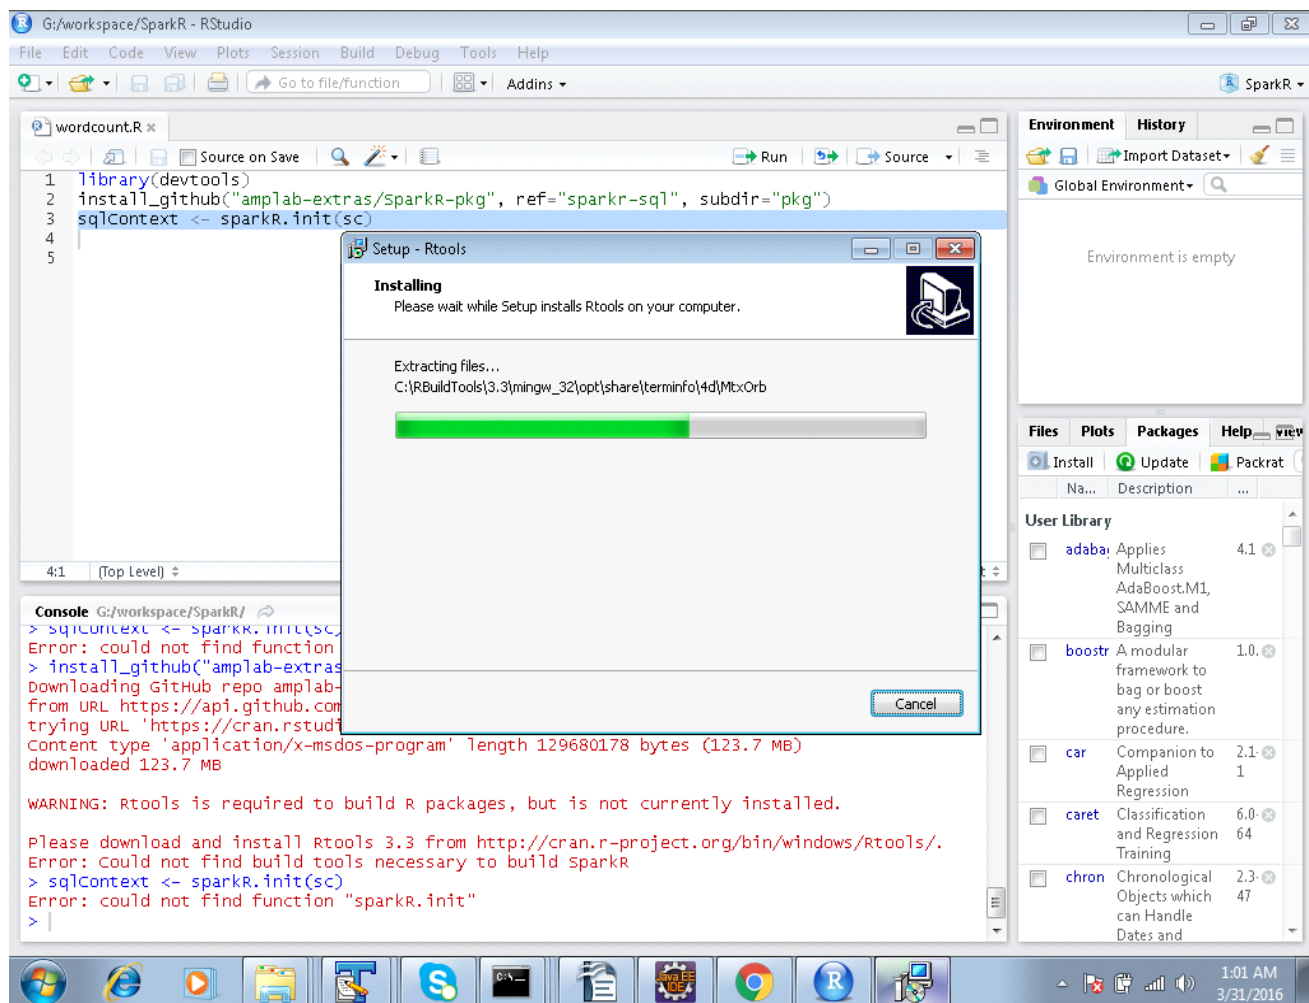
```
install_github("amplab-extras/SparkR-pkg", ref="sparkr-sql", subdir="pkg")
```

below are packages gets installed:

C:\RBuildTools\3.3

C:\RBuildTools\3.3\gcc-4.6.3

C:\RBuildTools\3.3\mingw_64



procedure 2:

```
library(devtools)
```

```
load_all("C:/spark-1.6.0-bin-hadoop2.3/R/lib/SparkR/R")
```

```
install("C:/spark-1.6.0-bin-hadoop2.3/R/lib/SparkR/R")
```

```
library(SparkR)  
sqlContext <- SparkR.init(sc)
```

install roxygen2 -since load_all requires roxygen2

Skipping missing files: schema.R, generics.R, jobj.R, RDD.R, pairRDD.R, column.R, group.R, DataFrame.R, SQLContext.R, backend.R, broadcast.R, client.R, context.R, deserialize.R, functions.R, mllib.R, serialize.R, sparkR.R, stats.R, types.R, utils.R

```
> install("C:/spark-1.6.0-bin-hadoop2.3/R/lib/SparkR/R")
```

```
Installing SparkR  
"C:/PROGRA~1/R/R-32~1.3/bin/x64/R" --no-site-file --no-environ --no-save --no-  
restore \  
  CMD INSTALL "C:/spark-1.6.0-bin-hadoop2.3/R/lib/SparkR" \  
  --library="C:/Users/Arun/Documents/R/win-library/3.2" --install-tests
```

procedure 3:

```
Sys.setenv(SPARK_HOME = "C:/spark-1.6.0-bin-hadoop2.3")
```

```
.libPaths(c(file.path(Sys.getenv("SPARK_HOME"), "R", "lib"), .libPaths()))
```

```
library(SparkR)
```

```
sc <- sparkR.init(master = "local")
```

Cmd to start sparkR:

```
C:\Windows\system32>sparkR
```



```
16/03/31 02:23:17 INFO util.Utils: Successfully started service 'SparkUI' on port 4041.
```

```
16/03/31 02:23:17 INFO ui.SparkUI: Started SparkUI at http://localhost:4041
```

```
16/03/31 02:23:17 INFO executor.Executor: Starting executor ID driver on host localhost
```

```
16/03/31 02:23:17 INFO util.Utils: Successfully started service 'org.apache.spark.network.netty.NettyBlockTransferService' on port 51230.
```

```
16/03/31 02:23:17 INFO netty.NettyBlockTransferService: Server created on 51230
```

```
16/03/31 02:23:17 INFO storage.BlockManagerMaster: Trying to register BlockManager
```

```
16/03/31 02:23:17 INFO storage.BlockManagerMasterEndpoint: Registering block manager localhost:51230 with 511.5 MB RAM, BlockManagerId(driver, localhost, 51230)
```

```
16/03/31 02:23:17 INFO storage.BlockManagerMaster: Registered BlockManager
```

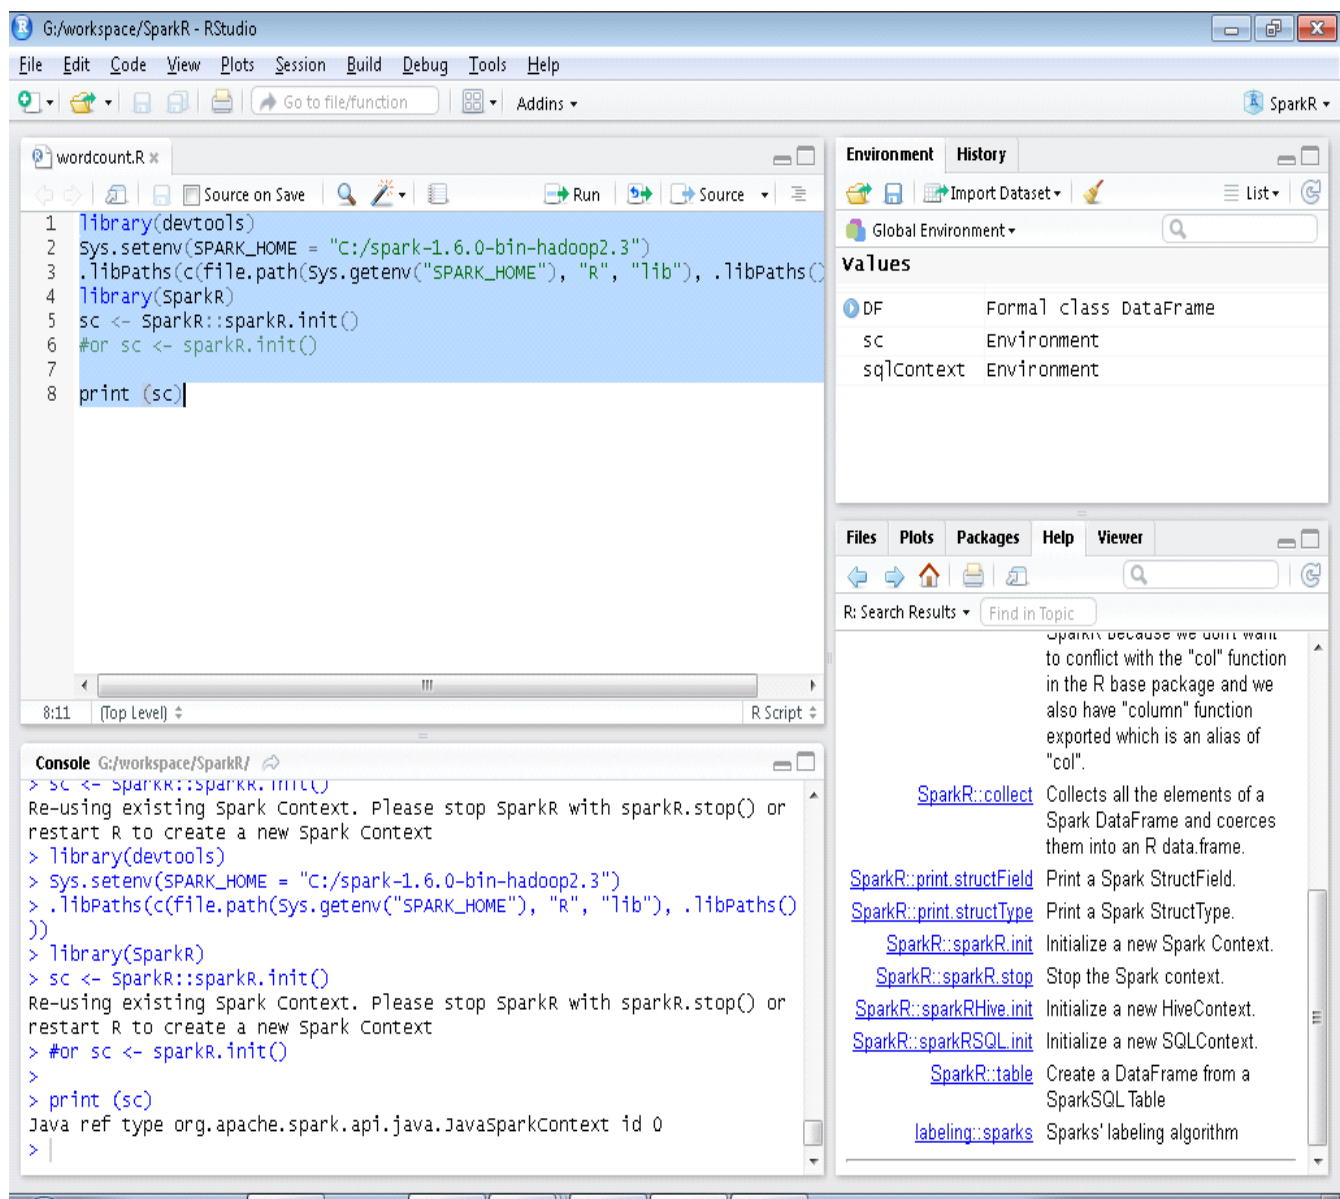
Welcome to



version 1.6.0

Spark context is available as `sc`, SQL context is available as `sqlContext`

>



Re-using existing Spark Context. Please stop SparkR with `sparkR.stop()` or restart R to create a new Spark Context

```
> read.df(sqlContext = sc, path = "hdfs://localhost:9000/input/wordcount.txt" )
```

Error in `invokeJava(isStatic = TRUE, className, methodName, ...)` :

org.apache.spark.SparkException: Job aborted due to stage failure: Task 3 in stage 9.0 failed 1 times, most recent failure: Lost task 3.0 in stage 9.0 (TID 39, localhost): java.io.IOException: Could not read footer: java.lang.RuntimeException: **hdfs://localhost:9000/input/wordcount.txt is not a Parquet file**. expected magic number at tail [80, 65, 82, 49] but found [109, 112, 108, 101]

Reading and writing

Advantages of R:

- R is highly dynamic
- easy data manipulation using data frames.
- powerful visualization

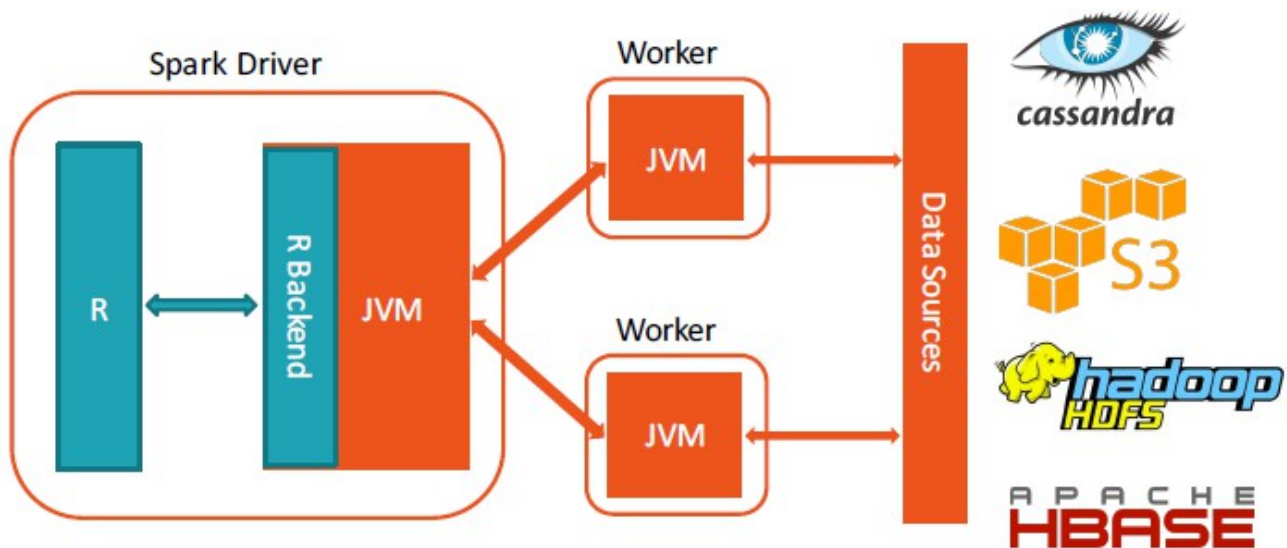
disadvantages of R:

- disadvantage is same -dynamic language features. I.e at runtime the loop variable cannot assume same datatype, every time it needs to check datatype.
- Single threaded.
- Everything has to fit in to memory.

SPARKR:

- provide R front end to spark

OVERVIEW OF SPARK



Reading and writing to storage(JVM <->Storage)

```
1.SPARKDF ← read.df(SqlContext,path = " _ " , source ="csv")
```

- jsonFile:read.df(sqlContext,path=" _ " ,source="json")
- parquetFile: read.df(sqlcontext , =" _ " ,source ="parquet")

NOTE:Above read from many source like database(e.g casandra),csv,json,

```
2.write.df (SparkDF,source="json")
```

```
saveAsparquetFile: write.df( _ , source="parquet")
```

NOTE:Above write from work to distributed storage and distributed storage to worker. Write uses df not sql context

NOTE:

As stated above **read data** from **distributed storage** to **worker** and **write** from **worker** to **distributed storage**.

To communicate with R process you need to follow below:

Moving data between R and JVM

only below commands talks with R process and jvm

- sparkDF <-createDataFrame(sqlcontext,df)
- df<- collect(sparkDF)

NOTE:Above read from many source

Caching

controls caching of distributed data

- persist(sparkDF,storagelevel)
- cache(sparkDF)
- cacheTable(sqlcontext,"tableName")


```
library(devtools)
library(SparkR)
Sys.setenv(SPARK_HOME = "C:/spark-1.6.0-bin-hadoop2.3")
.libPaths(c(file.path(Sys.getenv("SPARK_HOME"), "R", "lib"), .libPaths()))
Sys.getenv("SPARK_HOME")
#load the Sparkr library
library(SparkR)
# Create a spark context and a SQL context
sc <- sparkR.init(master = "local")
sqlContext <- sparkRSQL.init(sc)
sc <- sparkR.init(sparkPackages="com.databricks:spark-csv_2.11:1.0.3")
result <- read.df(sqlContext, "/input/sales.csv", "csv")
```

```
> result <- read.df(sqlContext, "/input/sales.csv", "csv")
Error in invokeJava(isStatic = TRUE, className, methodName, ...) :
  java.lang.ClassNotFoundException: Failed to find data source: csv. Please find
packages at http://spark-packages.org
```

install com.databricks:spark-csv <<version>> as needed in below command

```
> $SPARK_HOME/bin/spark-shell --packages com.databricks:spark-csv_2.11:1.0.3
```

see above command `sc <- sparkR.init(sparkPackages="com.databricks:spark-csv_2.11:1.0.3")`

is same as command prompt installation of spark-shell with `--packages` option.

The screenshot shows a web browser with the URL `spark-packages.org/package/databricks/spark-csv`. The page features the Spark Packages logo and navigation links: Feedback, Register a package, Login, and Find a package. The main content area displays the package `spark-csv` (homepage) as a Spark SQL CSV data source, credited to `@databricks` with a 5-star rating and 7 users. A description states: "This packages implements a CSV data source for Apache Spark. CSV files can be read as DataFrame." Below this, a 'Tags' section lists `sql`, `SparkSQL`, `DataSource`, and `csv`. The 'How to [+]' section instructs users to include the package in their Spark Applications using `spark-shell`, `pyspark`, or `spark-submit`, and provides a terminal command: `$SPARK_HOME/bin/spark-shell --packages com.databricks:spark-csv_2.11:1.4.0`. The 'Releases' section lists two versions: `1.4.0-s_2.11` and `1.4.0-s_2.10`, both dated 2016-03-05, with Apache-2.0 license and Scala versions 2.11 and 2.10 respectively. A footer note mentions that Spark Packages is a community site and that Apache Spark and its logo are trademarks of the Apache Software Foundation.

NOTE:

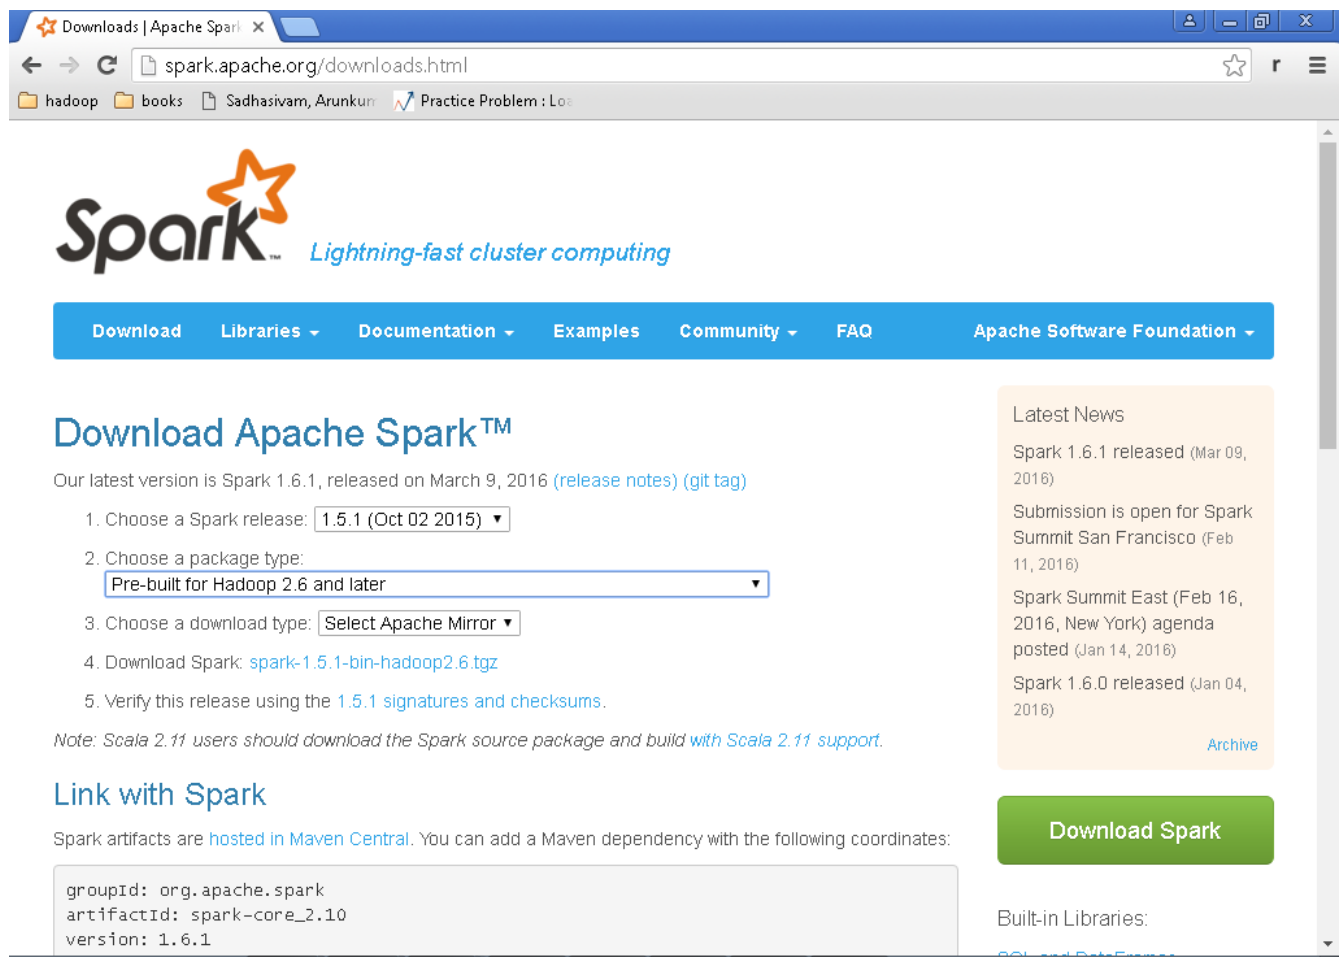
It might be worth mentioning that running `spark-1.5.1-bin-hadoop2.6/bin/spark-shell --packages com.databricks:spark-csv_2.11:1.2.0` works just fine.

Re-install with hadoop 2.6 and spark 1.5.1

STEP 1:

change the environment variable to point to 1.5.1 and close all cmd prompt

```
%SystemRoot%\system32;%SystemRoot%;%SystemRoot%\System32\Wbem;%SYSTEMROOT%
%\System32\WindowsPowerShell\v1.0\;C:\Program Files\Intel\WiFi\bin\;C:\Program Files\Common
Files\Intel\WirelessCommon\;C:\Program Files (x86)\Skype\Phone\;C:\apache-maven-
3.3.9\bin;C:\protoc;C:\Program Files\Microsoft SDKs\Windows\v7.1\bin;C:\Program
Files\Git\bin;C:\Java\jdk1.7.0_79\bin;C:\Anaconda2;C:\Anaconda2\Library\bin;C:\Anaconda2\Scripts;C:\Progra
m Files\R\R-3.2.3\bin;C:\spark-1.5.1-bin-hadoop2.6\bin;C:\scala-2.11.7\bin;C:\hadoop-2.2.0\bin;C:\hadoop-
2.2.0\sbin;C:\apache-mahout-distribution-0.10.2\bin;C:\pig-0.15.0\bin;C:\apache-hive-2.0.0-
bin\bin;C:\zeppelin-0.5.5\bin
```



The screenshot shows the Apache Spark download page in a web browser. The browser's address bar displays 'spark.apache.org/downloads.html'. The page features the Spark logo with the tagline 'Lightning-fast cluster computing'. A navigation bar includes links for 'Download', 'Libraries', 'Documentation', 'Examples', 'Community', 'FAQ', and 'Apache Software Foundation'. The main heading is 'Download Apache Spark™'. Below it, a message states: 'Our latest version is Spark 1.6.1, released on March 9, 2016 (release notes) (git tag)'. A five-step guide for downloading is provided: 1. Choose a Spark release (1.5.1 (Oct 02 2015) is selected), 2. Choose a package type (Pre-built for Hadoop 2.6 and later is selected), 3. Choose a download type (Select Apache Mirror is selected), 4. Download Spark (spark-1.5.1-bin-hadoop2.6.tgz is the link), and 5. Verify this release using the 1.5.1 signatures and checksums. A note mentions that Scala 2.11 users should download the source package and build with Scala 2.11 support. On the right, a 'Latest News' section lists recent releases and events. A large green 'Download Spark' button is prominent. At the bottom left, a code block shows Maven coordinates: groupId: org.apache.spark, artifactId: spark-core_2.10, version: 1.6.1. The bottom right corner mentions 'Built-in Libraries:'.

Downloads | Apache Spark

spark.apache.org/downloads.html

hadoop books Sadhasivam, Arunkun Practice Problem : Lo

Spark™

Lightning-fast cluster computing

Download Libraries Documentation Examples Community FAQ Apache Software Foundation

Download Apache Spark™

Our latest version is Spark 1.6.1, released on March 9, 2016 (release notes) (git tag)

1. Choose a Spark release: 1.5.1 (Oct 02 2015)
2. Choose a package type: Pre-built for Hadoop 2.6 and later
3. Choose a download type: Select Apache Mirror
4. Download Spark: spark-1.5.1-bin-hadoop2.6.tgz
5. Verify this release using the 1.5.1 signatures and checksums.

Note: Scala 2.11 users should download the Spark source package and build with Scala 2.11 support.

Link with Spark

Spark artifacts are hosted in Maven Central. You can add a Maven dependency with the following coordinates:

```
groupId: org.apache.spark
artifactId: spark-core_2.10
version: 1.6.1
```

Built-in Libraries:

Latest News

- Spark 1.6.1 released (Mar 09, 2016)
- Submission is open for Spark Summit San Francisco (Feb 11, 2016)
- Spark Summit East (Feb 16, 2016, New York) agenda posted (Jan 14, 2016)
- Spark 1.6.0 released (Jan 04, 2016)

Archive

Download Spark

STEP 2:

Try running java , see below it stills uses 1.6.1.

```
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
16/04/12 22:36:28 INFO SparkContext: Running Spark version 1.6.1
16/04/12 22:36:29 INFO SecurityManager: Changing view acls to: Arun
16/04/12 22:36:29 INFO SecurityManager: Changing modify acls to: Arun
16/04/12 22:36:29 INFO SecurityManager: SecurityManager: authentication disabled; ui acls
disabled; users with view permissions: Set(Arun); users with modify permissions: Set(Arun)
16/04/12 22:36:30 INFO Utils: Successfully started service 'sparkDriver' on port 58098.
16/04/12 22:36:31 INFO Slf4jLogger: Slf4jLogger started
16/04/12 22:36:31 INFO Remoting: Starting remoting
```

change the below env entry-SPARK_JAR

C:\spark-1.5.1-bin-hadoop2.6\lib\spark-assembly-1.5.1-hadoop2.6.0.jar

change pom.xml

<dependencies>

<dependency>

<groupId>org.apache.spark</groupId>

<artifactId>spark-streaming_2.10</artifactId>

<version>1.5.1</version>

</dependency>

</dependencies>

check java output and version

```
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
16/04/12 22:45:34 INFO SparkContext: Running Spark version 1.5.1
16/04/12 22:45:35 INFO SecurityManager: Changing view acls to: Arun
16/04/12 22:45:35 INFO SecurityManager: Changing modify acls to: Arun
16/04/12 22:45:35 INFO SecurityManager: SecurityManager: authentication disabled; ui acls
disabled; users with view permissions: Set(Arun); users with modify permissions: Set(Arun)
16/04/12 22:45:36 INFO Slf4jLogger: Slf4jLogger started
16/04/12 22:45:36 INFO Remoting: Starting remoting
16/04/12 22:45:36 INFO Remoting: Remoting started; listening on addresses :
[akka.tcp://sparkDriver@localhost:58361]
16/04/12 22:45:36 INFO Utils: Successfully started service 'sparkDriver' on port 58361.
16/04/12 22:45:41 INFO DAGScheduler: Job 0 finished: collect at JavaWordCountHdfs.java:76,
took 0.585637 s
```

example: 1

are: 1

```
is: 1
you: 1
wordcount: 1
hadoop: 3
hi: 3
how: 2
16/04/12 22:45:41 INFO SparkUI: Stopped Spark web UI at http://localhost:4040
16/04/12 22:45:41 INFO DAGScheduler: Stopping DAGScheduler
16/04/12 22:45:41 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint
stopped!
```

Now try running sparkR R program to Access HDFS:

```
c:\spark-1.5.1-bin-hadoop2.6\bin> spark-shell --packages com.databricks:spark-csv_2.11:1.2.0
```

Note: on running the above command no error is showing.

```
16/04/12 22:54:11 INFO ui.SparkUI: Started SparkUI at http://localhost:4040
16/04/12 22:54:11 INFO spark.SparkContext: Added JAR file:/C:/Users/Arun/.ivy2/j
ars/com.databricks_spark-csv_2.11-1.2.0.jar at http://localhost:54649/jars/com.d
atabricks_spark-csv_2.11-1.2.0.jar with timestamp 1460481851695
16/04/12 22:54:11 INFO spark.SparkContext: Added JAR file:/C:/Users/Arun/.ivy2/j
ars/org.apache.commons_commons-csv-1.1.jar at http://localhost:54649/jars/org.ap
ache.commons_commons-csv-1.1.jar with timestamp 1460481851701
16/04/12 22:54:11 INFO spark.SparkContext: Added JAR file:/C:/Users/Arun/.ivy2/j
ars/com.univocity_univocity-parsers-1.5.1.jar at http://localhost:54649/jars/com
.univocity_univocity-parsers-1.5.1.jar with timestamp 1460481851706
16/04/12 22:54:11 WARN metrics.MetricsSystem: Using default name DAGScheduler fo
r source because spark.app.id is not set.
```


NOTE:

See all jar are automatically deployed to sucessfully when running below command

```
spark-shell --packages com.databricks:spark-csv_2.11:1.2.0
```

both jar spark-csv_2.11-1.2.0 downloaded should match the jar add to class path in SparkR programming.

Add spark csv (databricks) to class path of R programming.



```

library(devtools)

library(SparkR)

Sys.setenv(SPARK_HOME = "C:/spark-1.5.1-bin-hadoop2.6")

.libPaths(c(file.path(Sys.getenv("SPARK_HOME"), "R", "lib"), .libPaths()))

Sys.getenv("SPARK_HOME")

#load the Sparkr library

library(SparkR)

# Create a spark context and a SQL context

sc <- sparkR.init(master = "local")

sqlContext <- sparkRSQL.init(sc)

sc <- sparkR.init(sparkPackages="com.databricks:spark-csv_2.11:1.2.0")

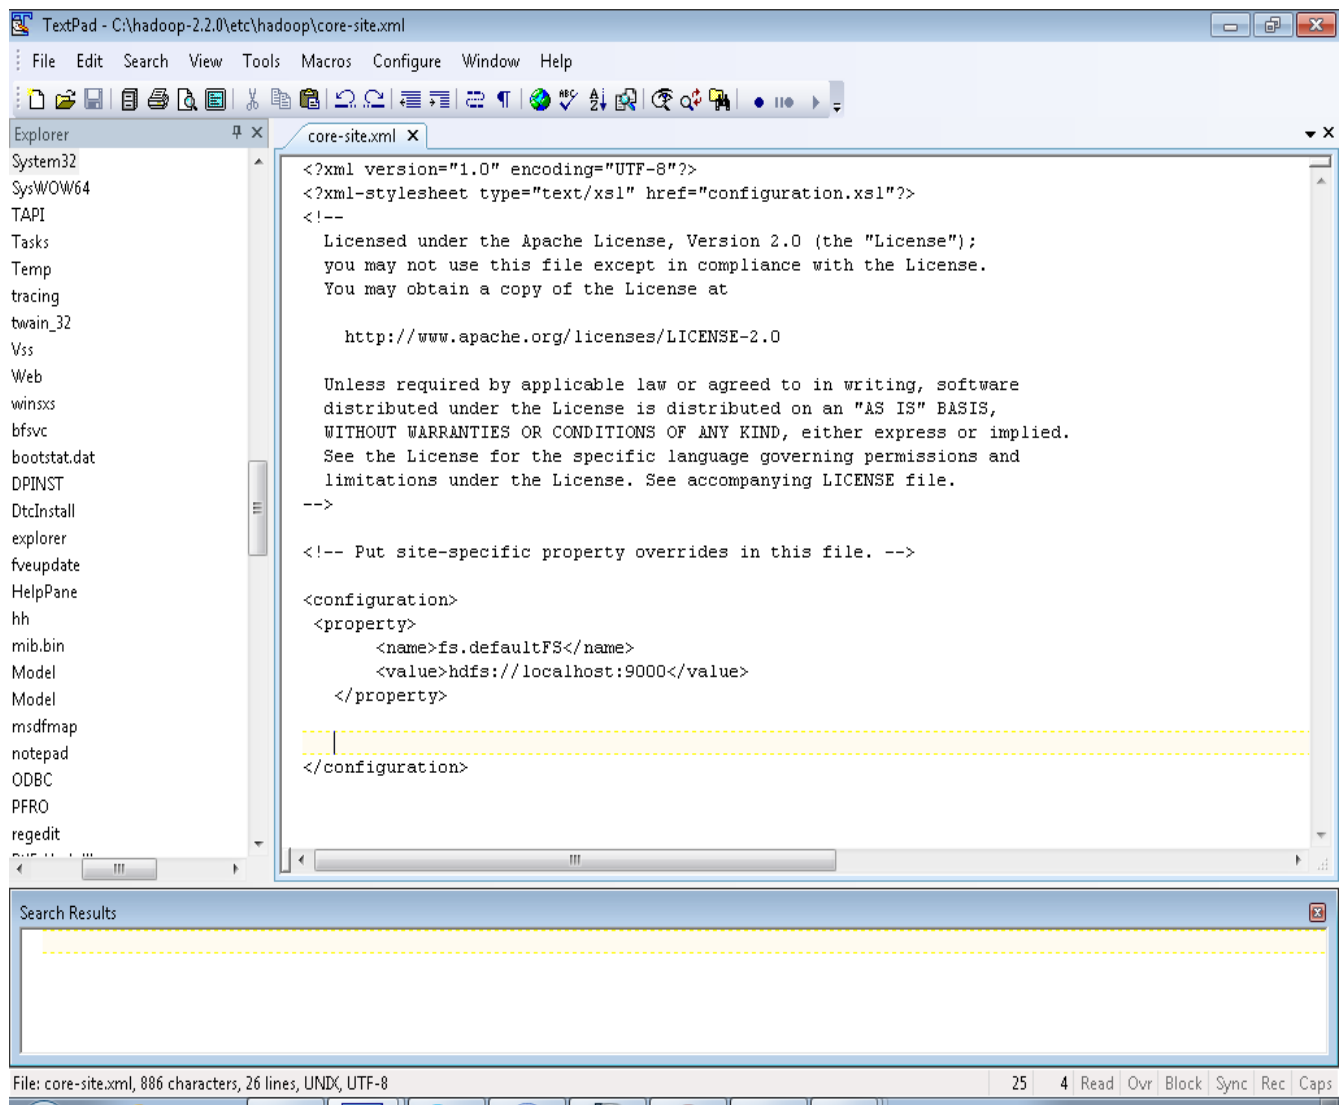
result <- read.df(sqlContext, "hdfs://localhost:9000/input/sales.csv")

```

```

Error in invokeJava(isStatic = TRUE, className, methodName, ...) :
  org.apache.spark.SparkException: Job aborted due to stage failure: Task 0 in
stage 0.0 failed 1 times, most recent failure: Lost task 0.0 in stage 0.0 (TID 0,
localhost): java.io.IOException: Could not read footer:
java.lang.UnsatisfiedLinkError:
org.apache.hadoop.util.NativeCrc32.nativeComputeChunkedSums(IILjava/nio/ByteBuffer
;ILjava/nio/ByteBuffer;IILjava/lang/String;JZ)V
    at
org.apache.parquet.hadoop.ParquetFileReader.readAllFootersInParallel(ParquetFileRe
ader.java:247)
    at org.apache.spark.sql.execution.datasources.parquet.ParquetRelation$
$anonfun$28.apply(ParquetRelation.scala:750)
    at org.apache.spark.sql.execution.datasources.parquet.ParquetRelation$
$anonfun$28.apply(ParquetRelation.scala:739)
    at org.apache.spark.rdd.RDD$$anonfun$mapPartitions$1$
$anonfun$apply$17.apply(RDD.scala:706)
    at org.apache.spark.rdd.RDD$$anonfun$mapPartitions$1$
$anonfun$apply$17.apply(RDD.scala:706)
    at
org.apache.spark.rdd.MapPartitionsRDD.compute(MapPartitionsRDD.scala:38)

```



NOTE:

hdfs root path in R should be same as given in core-site.xml

e.g result <- read.df(sqlContext, "hdfs://localhost:9000/input/sales.csv")

